

DATA403 Final Presentation:

with HumanoidStandup-v4
and HalfCheetah-v4


2021320303
Jiwon Jeong

Agenda

- Introduction
- Method
- Experiments & Analysis (HumanoidStandup-v4)
- Experiments & Analysis (HalfCheetah-v4)
- Conclusion


Introduction

- HumanoidStandup-v4
 - Goal: Make the humanoid standup and then keep it standing
 - Action Space: 17 continuous actions in $[-0.4, 0.4]$
 - Action represents the numerical torques applied at the hinge joints
 - Observation Space: 376 continuous space in $(-\infty, +\infty)$
 - State consists of positional values of different body parts of Humanoid
 - Why? Very Challenging and Interesting (Unsolved)




Introduction

- HalfCheetah-v4
 - Goal: Make the cheetah run forward as fast as possible
 - Action Space: 6 continuous actions in $[-1, 1]$
 - Action represents the torques applied between links
 - Observation Space: 17 continuous space in $(-\infty, +\infty)$
 - State consists of positional values of different body parts of the cheetah
 - Why? Looks Fun !




Method


- Based on PPO (Proximal Policy Gradient)
- MLP with tanh




State → 64 → 64 → Action



State → 128 → 64 → Action



State → 64 → 32 → Action



State → 128 → 64 → 64 → Action

etc...

Method

- Gym Wrappers
 - FlattenObservation – flatten the observation
 - ClipAction – clip the continuous action to the valid bound
 - NormalizeObservation – normalize observation
 - TransformObservation – clip(obs, -10, 10)
 - NormalizeReward – normalize reward
 - TransformReward – clip(reward, -10, 10)

Method

- Robust Policy Optimization (ICLR 2023)

ROBUST POLICY OPTIMIZATION IN DEEP REINFORCE-
MENT LEARNING

Md Masudur Rahman & Yexiang Xue
Department of Computer Science
Purdue University
West Lafayette, IN 47907, USA
{rahman64, yexiang}@purdue.edu

- Modified from PPO
- RPO leverages a method of perturbing the distribution representing actions
- Improved Performance compared to PPO

Method

- Robust Policy Optimization (ICLR 2023)




Algorithm 1 Robust Policy Optimization (RPO)

```
1: Initialize parameter vectors  $\theta$  for policy network.
2: for each iteration do
3:    $\mathcal{D} \leftarrow \{\}$ 
4:   for each environment step do
5:      $\mu, \sigma \leftarrow \pi_{\theta}(\cdot | s_t)$ 
6:      $a_t \sim \mathcal{N}(\mu, \sigma)$ 
7:      $s_{t+1} \sim P(s_{t+1} | s_t, a_t)$ 
8:      $r_t \sim R(s_t, a_t)$ 
9:      $\mathcal{D} \leftarrow \mathcal{D} \cup \{(s_t, a_t, r_t, s_{t+1})\}$ 
10:  end for
11:  for each observation  $s_t$  in  $\mathcal{D}$  do
12:     $\mu, \sigma \leftarrow \pi_{\theta}(\cdot | s_t)$ 
13:     $z \sim \mathcal{U}(-\alpha, \alpha)$ 
14:     $\mu' \leftarrow \mu + z$ 
15:     $prob \leftarrow \mathcal{N}(\mu', \sigma)$ 
16:     $logp \leftarrow prob(a_t)$ 
17:    Compute RL loss  $L_{\pi}$  using  $logp$ ,  $a_t$ , and value function.
18:  end for
19: end for
```

```
def get_action_and_value(self, x, action=None):
    ##### Implement here : policy distribution #####
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
    action_mean = self.actor_mean(x)
    action_logstd = self.actor_logstd.expand_as(action_mean)
    action_std = torch.exp(action_logstd)
    probs = Normal(action_mean, action_std)
    if action is None:
        action = probs.sample()
    else: # RPO
        z = torch.FloatTensor(action_mean.shape).uniform_(-self.rpo_alpha, self.rpo_alpha).to(device)
        action_mean = action_mean + z
        probs = Normal(action_mean, action_std)
```


Method

- “Symmetric Action Loss” in HumanoidStandup setting
 - Why do Humanoid use only one arm or leg?



Method

- “Symmetric Action Loss” in HumanoidStandup setting
 - Gym Documentation: Action descriptions

6	Torque applied on the rotor between the right hip/thigh and the right shin	-0.4	0.4	right_knee	hinge	torque (N m)
10	Torque applied on the rotor between the left hip/thigh and the left shin	-0.4	0.4	left_knee	hinge	torque (N m)

- Implementation (Use MSELoss() in torch.nn)

```
##### Symmetric action loss #####
sym_act_loss_1 = self.action_loss(probs.log_prob(action)[: , 3:7], probs.log_prob(action)[: , 7:11])
sym_act_loss_2 = self.action_loss(probs.log_prob(action)[: , 11:14], probs.log_prob(action)[: , 14:17])
sym_act_loss = 0.5 * (sym_act_loss_1 + sym_act_loss_2)
```

Experiments

- Device

- CPU: AMD Ryzen 5 3600
- GPU: GTX 1060 3GB




- Changed Hyperparameters

- Seed
- Entropy coefficient, Symmetric action coefficient
- Learning rate
- RPO alpha

Experiments (HumanoidStandup-v4)


- 1. Entropy Coefficient
 - 0.01 (pink) vs. **0.0001 (yellow)** vs. 0.0 (blue)



Tensorboard


Experiments (HumanoidStandup-v4)

- 2. Learning rate
 - $3e-4$ (gray) with 1M timesteps vs. $1e-4$ (cyan) with 2M timesteps



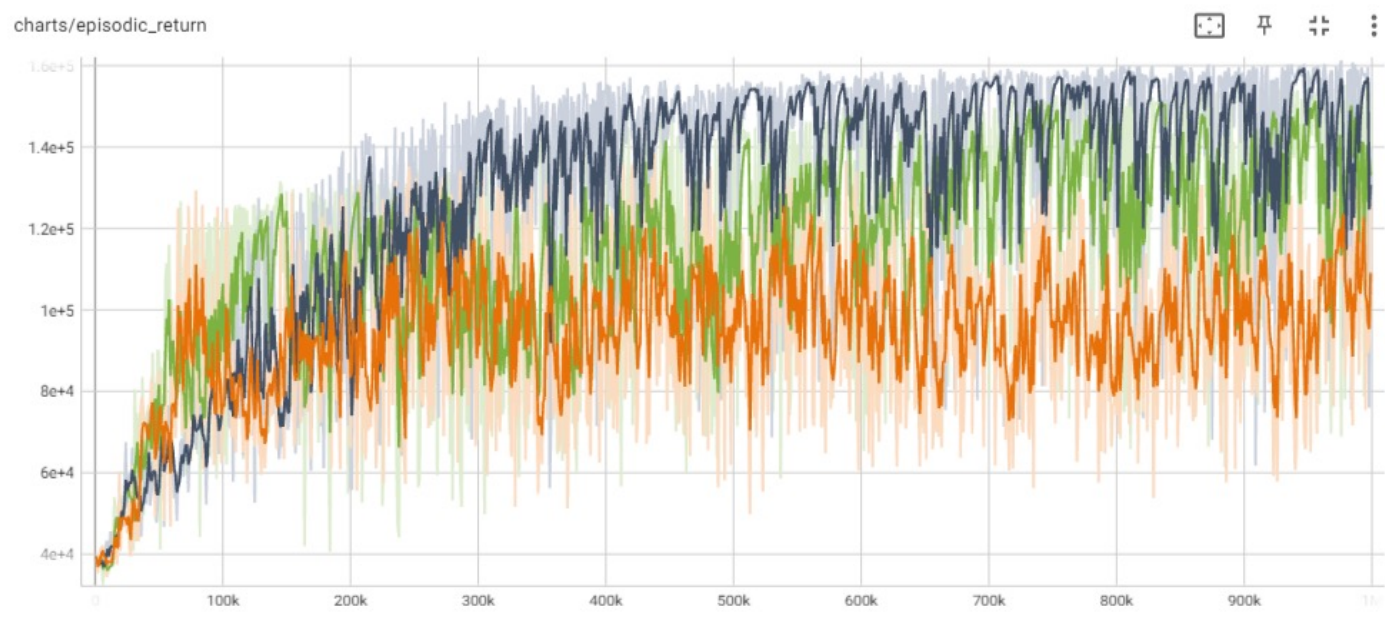
Experiments (HumanoidStandup-v4)

- 3. PPO vs. RPO
 - PPO (green) vs. **RPO with coef=0.5 (yellow)**



Experiments (HumanoidStandup-v4)




- 4. Symmetric action coefficient
 - 0.0 (green) vs. **0.01 (black)** vs. 0.05 (orange). Positive effect!



Experiments (HumanoidStandup-v4)

- Best Result

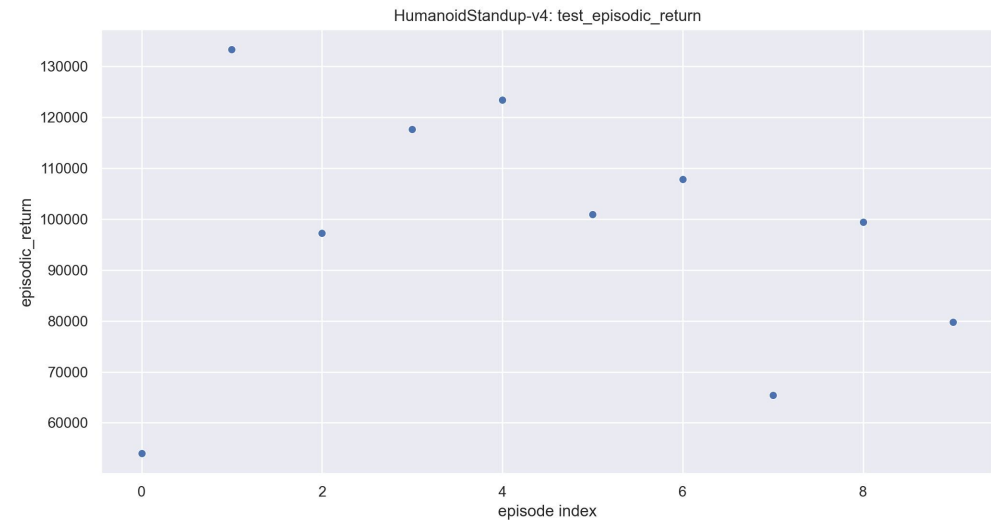
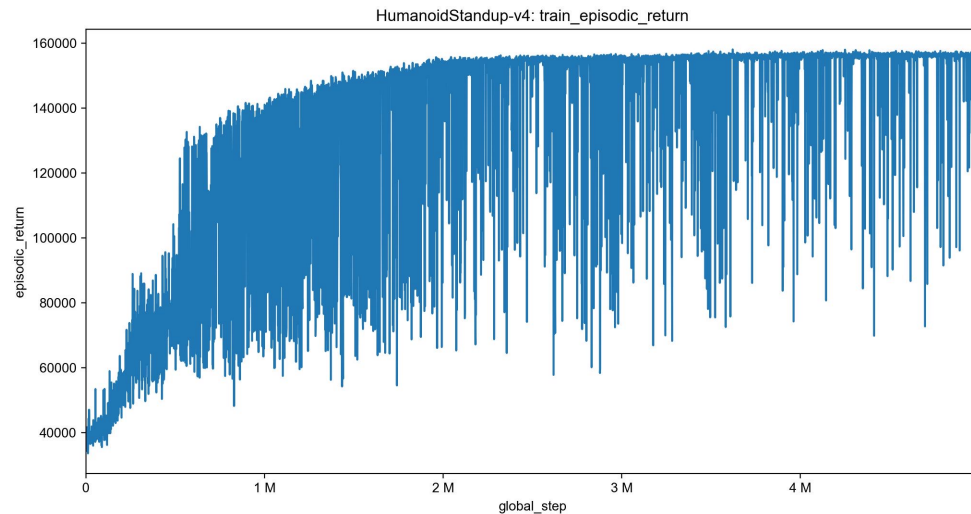
- Training: 160,000 (highest)
- Test: 131,000 (highest)



Hyperparameter	Value
Seed	0
Timesteps	5000000 (5M)
num_rollout_steps	2048
minibatch_size	64
Learning rate	2e-5
max_grad_norm	0.5
clip_coef	0.2
ent_coef	0.0001
vf_coef	0.5
gamma	0.99
gae_lambda	0.95
rpo_coef	0.5
sym_action_coef	0.02

Experiments (HumanoidStandup-v4)

- Best Result






Analysis (HumanoidStandup-v4)

- RPO is better than PPO

```
def get_action_and_value(self, x, action=None):
    ##### Implement here : policy distribution #####
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
    action_mean = self.actor_mean(x)
    action_logstd = self.actor_logstd.expand_as(action_mean)
    action_std = torch.exp(action_logstd)
    probs = Normal(action_mean, action_std)
    if action is None:
        action = probs.sample()
    else: # RPO
        z = torch.FloatTensor(action_mean.shape).uniform_(-self.rpo_alpha, self.rpo_alpha).to(device)
        action_mean = action_mean + z
        probs = Normal(action_mean, action_std)
```


Analysis (HumanoidStandup-v4)

- Symmetric Action Loss guided the use of both arms and legs
- And Improve performance!




Experiments (HalfCheetah-v4)

- 1. Entropy Coefficient
 - **0.0 (pink)** vs. 0.0001 (black)



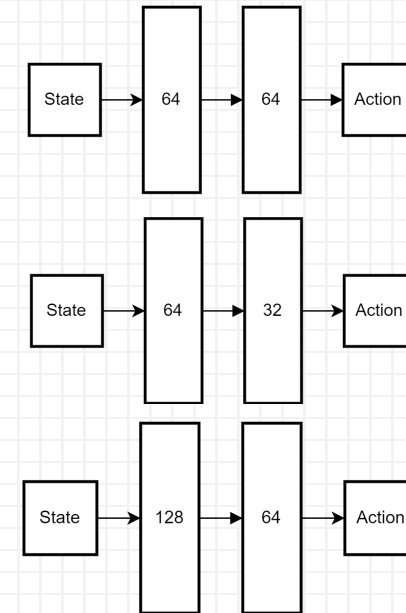

Experiments (HalfCheetah-v4)

- 2. RPO vs PPO
 - **PPO (orange)** vs. RPO with coef=0.5 (cyan)





Experiments (HalfCheetah-v4)

- 3. MLP
 - **64-64 (Orange)** vs. 64-32 (cyan) vs. 128-64 (yellow)



Experiments (HalfCheetah-v4)




- 4. Seed
 - 0 (left) vs. 8 (right)



Experiments (HalfCheetah-v4)

- Best Performance

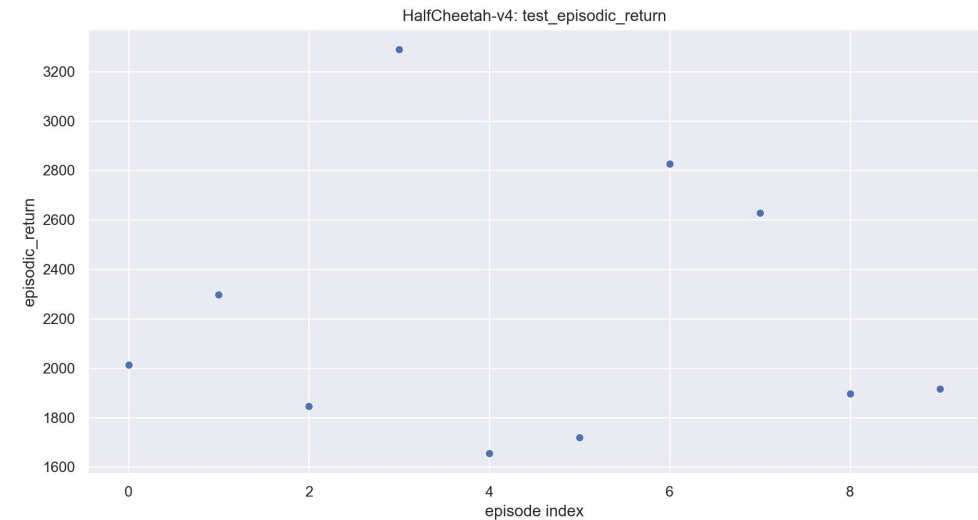
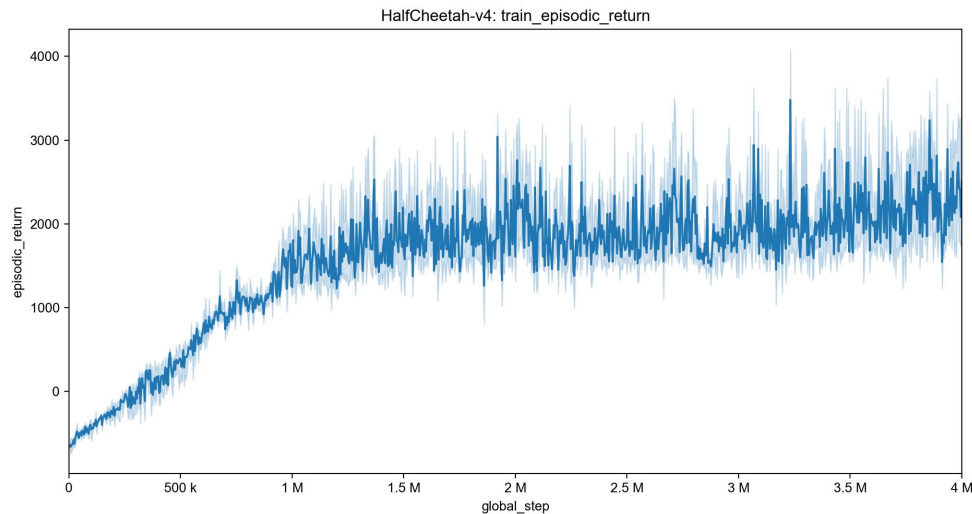
- Training: 4,000 (highest)
- Test: 3,250 (highest)



Hyperparameter	Value
Seed	8
Timesteps	1000000 (1M)
num_rollout_steps	128
minibatch_size	32
Learning rate	3e-5
max_grad_norm	2.0
clip_coef	0.2
ent_coef	0.5
vf_coef	0.5
gamma	0.99
gae_lambda	0.95

Experiments (HalfCheetah-v4)

- Best Result



Analysis (HalfCheetah-v4)


- Why does halfcheetah do a forward roll ?
 - Maybe halfcheetah get a big reward at once... → Stuck in local optima

Rewards

The reward consists of two parts:

- *forward_reward*: A reward of moving forward which is measured as $\text{forward_reward_weight} * (x\text{-coordinate before action} - x\text{-coordinate after action})/dt$. dt is the time between actions and is dependent on the `frame_skip` parameter (fixed to 5), where the `frametime` is 0.01 - making the default $dt = 5 * 0.01 = 0.05$. This reward would be positive if the cheetah runs forward (right).
- *ctrl_cost*: A cost for penalising the cheetah if it takes actions that are too large. It is measured as $\text{ctrl_cost_weight} * \text{sum}(\text{action}^2)$ where `ctrl_cost_weight` is a parameter set for the control and has a default value of 0.1


The total reward returned is **reward** = *forward_reward* - *ctrl_cost* and `info` will also contain the individual reward terms



Depending on the seed, Halfcheetah do or do not a forward roll

Analysis (HalfCheetah-v4)

- Why does halfcheetah do a forward roll ?
 - Maybe halfcheetah get a big reward at once... → Stuck in local optima
 - Forward roll: ~1500 (left) vs. Run fast: ~3000+ (right)



Conclusion

- Hyperparameters are **very important**
- Entropy loss is not useful in the two environments
- RPO worked depending on the environment
- Depending on the seed, results can vary a lot...
- It seems that other methods are needed to standup a humanoid.
- Still, PPO is an effective reinforcement learning algorithm!

Thanks 😊

